

RTS Project 1

Part I: Linux Kernel Building

Part II: Linux Scheduling Policy Analysis

Advisor: Prof. Tei-Wei Kuo

TA: Wen Sheng Lim
d08922028@csie.ntu.edu.tw


Slides credit: Han-Yi Lin

Outline

- ▶ What is “Kernel”
- ▶ Environment Setup
- ▶ Build Linux Kernel
- ▶ Linux Scheduling Policy Analysis
- ▶ Project Requirements
- ▶ Submission Rules

What is “Kernel”?

- ▶ The kernel^[1] is a fundamental part of a modern computer's operating system.
- ▶ The kernel's primary functions are to
 - ▶ Manage the computer's hardware and resources
 - ▶ E.g., CPU, main memory, I/O devices, and so on.
 - ▶ Allow applications to run and use these resources



Environment Setup

- ▶ Oracle VM VirtualBox^[2]
 - ▶ <https://www.virtualbox.org/wiki/Downloads>
- ▶ Ubuntu 20.04 64bits LTS (or newer) ^[3]
 - ▶ <https://ubuntu.com/download/desktop>
- ▶ Install the Ubuntu on the VirtualBox
 - ▶ Recommended: memory \geq 4GB, disk \geq 60GB, assign 1 CPU for VM to avoid complexity of multi-core scheduling

Build Linux Kernel (1 / 5)

- After the installation, please login Ubuntu and open a terminal to start building your Linux kernel^[4]

```
rts2022@rts2022-VirtualBox:~$ uname -a
Linux rts2022-VirtualBox 5.15.0-52-generic #58~20.04.1-Ubuntu SMP
Thu Oct 13 13:09:46 UTC 2022 x86_64 x86_64 x86_64 GNU/Linux
rts2022@rts2022-VirtualBox:~$
```

Build Linux Kernel (2/5)

- ▶ `$ sudo apt-get install build-essential libncurses5-dev flex bison libssl-dev libelf-dev dwarves zstd mpv`
- ▶ `$ wget`
<https://cdn.kernel.org/pub/linux/kernel/v5.x/linux-5.15.71.tar.xz>
- ▶ `$ tar xvf linux-5.15.71.tar.xz`
- ▶ `$ cd linux-5.15.71`
- ▶ `$ make mrproper`

Build Linux Kernel (3/5)

- ▶ `$ make menuconfig`
 - ▶ Make `CONFIG_SYSTEM_TRUSTED_KEYS` and `CONFIG_SYSTEM_REVOCATION_KEYS` empty
 - ▶ Disable `CONFIG_DEBUG_INFO_BTTF`
- ▶ `$ make`
 - ▶ You can use `make -j#` (# is the number of your physical cores) to create multiple threads to speed up the kernel building
- ▶ `$ sudo make modules_install`

Build Linux Kernel (4/5)

- ▶ `$ sudo make install`
- ▶ `$ sudoedit /etc/default/grub`
 - ▶ Configure the following 2 items
 - ▶ `GRUB_TIMEOUT_STYLE=menu`
 - ▶ `GRUB_TIMEOUT=10`
- ▶ `$ sudo update-grub`
- ▶ `$ sudo reboot`

Build Linux Kernel (5/5)

- Now, you can select the kernel you installed in the GNU grub to boot your Ubuntu.

```
GNU GRUB  version 2.04

*Ubuntu, 採用 Linux 5.15.71
Ubuntu, with Linux 5.15.71 (recovery mode)
Ubuntu, 採用 Linux 5.15.71.old
Ubuntu, with Linux 5.15.71.old (recovery mode)
Ubuntu, 採用 Linux 5.15.0-52-generic
Ubuntu, with Linux 5.15.0-52-generic (recovery mode)
Ubuntu, 採用 Linux 5.15.0-50-generic
Ubuntu, with Linux 5.15.0-50-generic (recovery mode)
```

- Then, you can use terminal and type “uname -a” to check the kernel version.

```
rts2022@rts2022-VirtualBox:~$ uname -a
Linux rts2022-VirtualBox 5.15.71 #4 SMP Fri Oct 21 23:34:30 CST
2022 x86_64 x86_64 x86_64 GNU/Linux
```

References

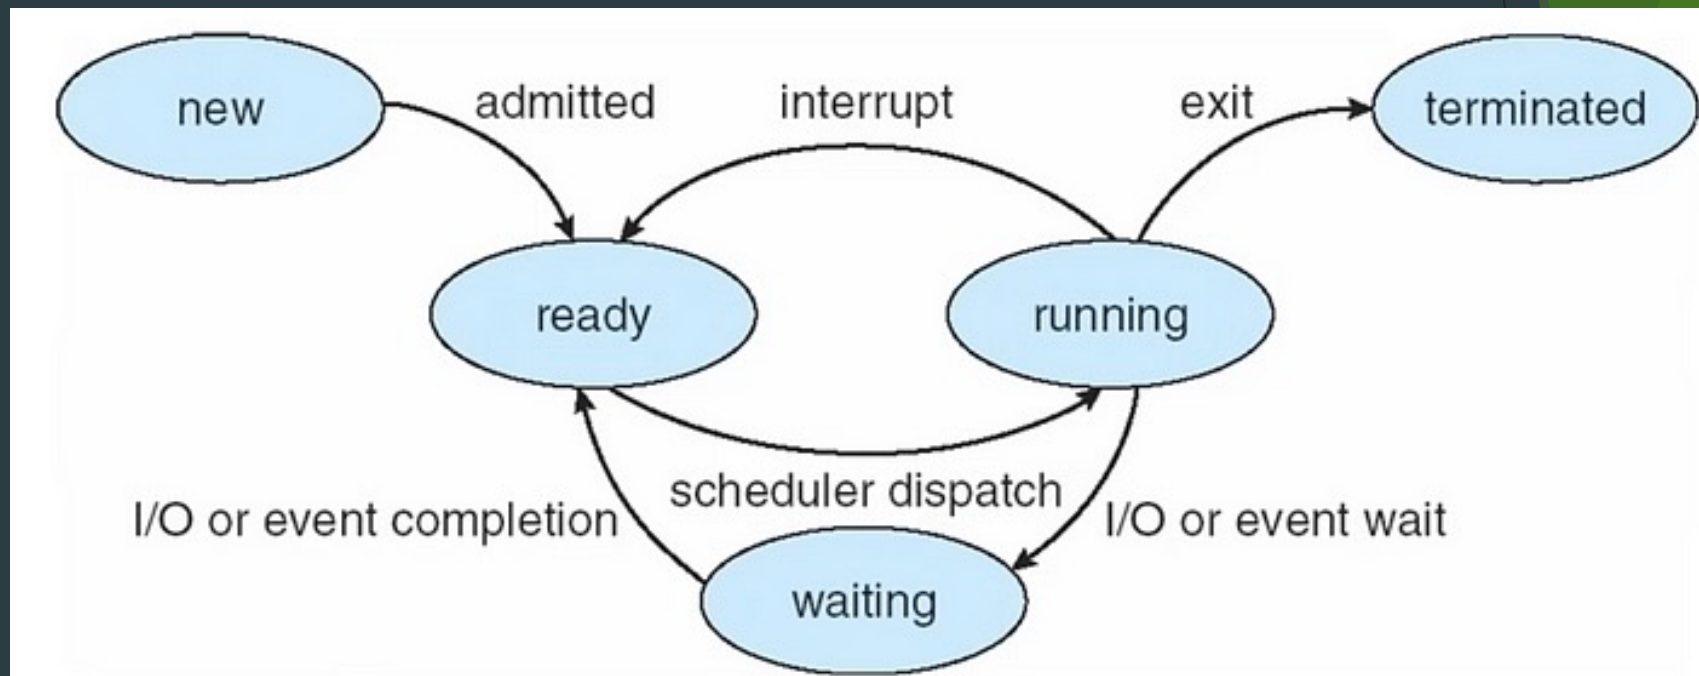
- ▶ [1] Wikipedia [https://en.wikipedia.org/wiki/Kernel_\(operating_system\)](https://en.wikipedia.org/wiki/Kernel_(operating_system))
- ▶ [2] Oracle VM VirtualBox <https://www.virtualbox.org/>
- ▶ [3] Ubuntu <https://ubuntu.com/>
- ▶ [4] Linux Kernel in a Nutshell <http://www.kroah.com/lkn/>

SCHEDULING IN LINUX (補充資料)

Process Life Cycle

- ▶ A process is not always ready to run.
- ▶ The scheduler must know the status of every process in the system when switching between tasks.
- ▶ A process may have one of the following **states**:
 - ▶ Running — The process is **executing at the moment**.
 - ▶ Waiting — The process is able to run but is not allowed to **because the CPU is allocated to another process**. The scheduler can select the process at the next task switch.
 - ▶ Sleeping — The process is sleeping and cannot run because it is **waiting for an external event**. The scheduler cannot select the process at the next task switch.
- ▶ The system saves all processes in a process table.

Transitions between Process States



The Need of the Scheduler

- ▶ A unique description of each process is held in memory and is linked with other processes by means of several structures.
- ▶ This is the situation facing the scheduler, whose task is *to share CPU time between the programs to create the illusion of concurrent execution*.
- ▶ Scheduler responsibilities —
 - ▶ Scheduling policy
 - ▶ Context switching

Linux Scheduling Subsystem

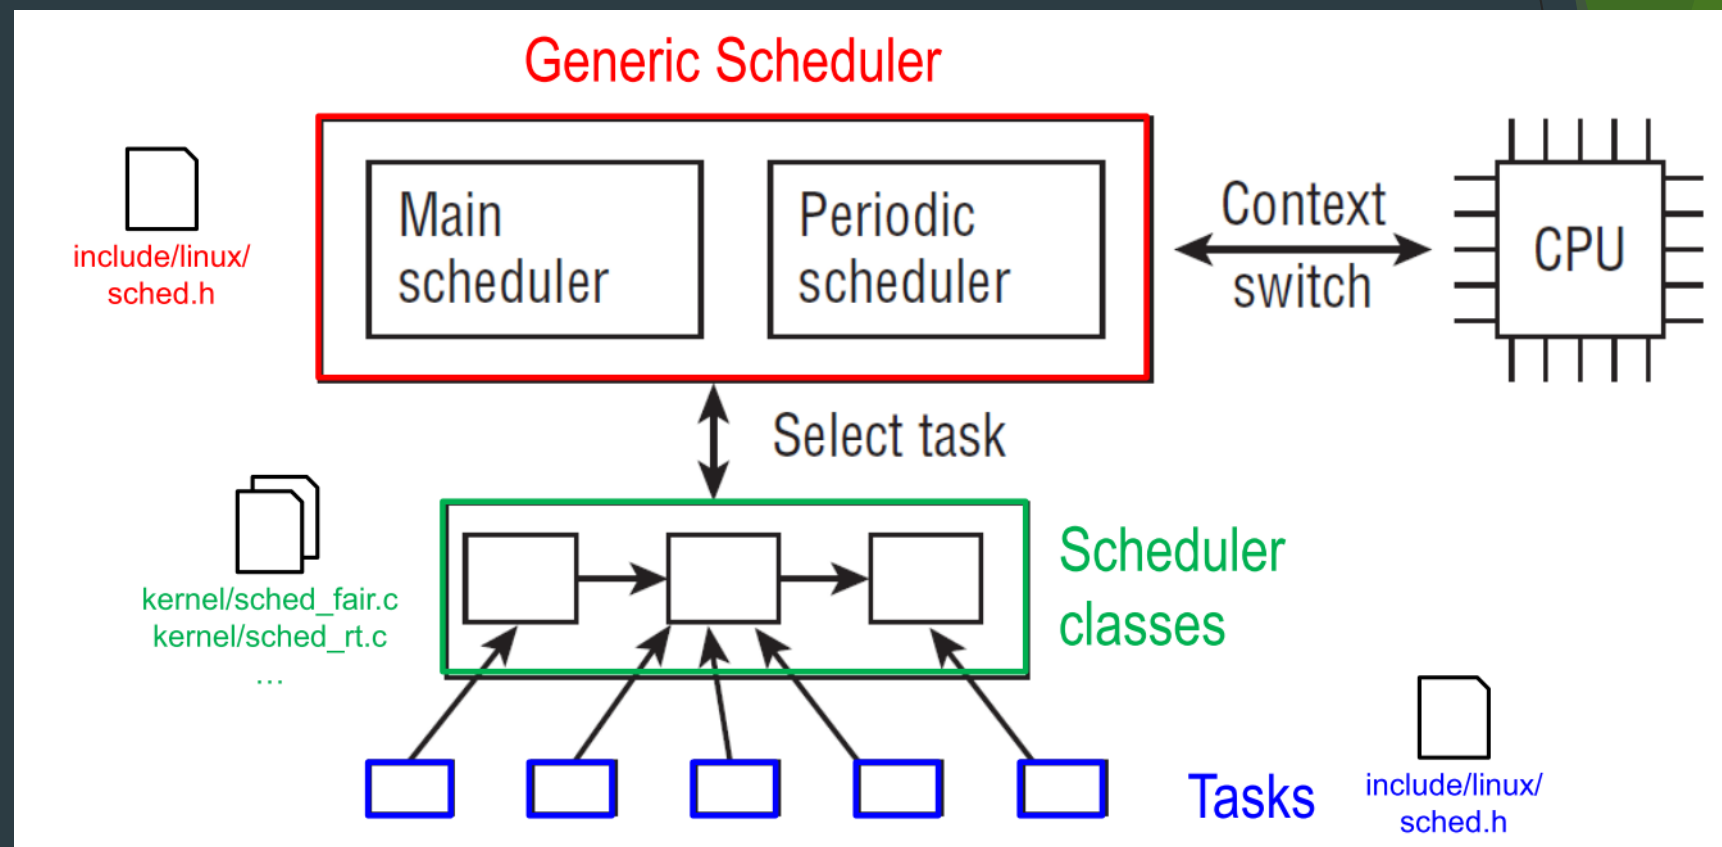
Generic Scheduler

Scheduler Classes

Task

Task

Task



Scheduling in Linux (1/2)

- ▶ The schedule function is the starting point to an understanding of scheduling operations.
- ▶ It is defined in “`kernel/sched/core.c`” and is one of the most frequently invoked functions in the kernel code.
- ▶ Not only priority scheduling but also two other soft real-time policies required by the POSIX standard are implemented.
- ▶ E.g., completely fair scheduling, real-time scheduling and scheduling of the idle task, etc.

Scheduling in Linux (2/2)

- ▶ The scheduler uses a series of data structures to sort and manage the processes in the system.
- ▶ Scheduling can be activated in two ways:
 - ▶ **Main scheduler**: Either directly if a task goes to sleep or wants to yield the CPU for other reasons,
 - ▶ **Periodic scheduler**: Or by a periodic mechanism that is run with constant frequency to check from time to time if switching tasks is necessary
- ▶ Generic scheduler = Main + Periodic schedulers




Linux Scheduling Policy Testing and Analysis

Linux Scheduling Policy Testing and Analysis

- ▶ Linux Scheduling Policy Classes
 - ▶ Normal Scheduling policies (Non-real-time)
 - ▶ SCHED_OTHER, SCHED_BATCH, SCHED_IDLE.
 - ▶ Real-Time policies
 - ▶ SCHED_FIFO, SCHED_RR, SCHED_DEADLINE.
- ▶ The default scheduling policy is **non-real-time**.
- ▶ In this part, using Linux **deadline scheduling policy (EDF)** to schedule threads in a process.

The Linux Deadline Scheduler

- ▶ Earliest Deadline First (EDF)
 - ▶ Each task is tracked as a node of a red-black tree
 - ▶ Nodes are ordered by dynamic scheduling deadline



Using Real-time Schedulers

- ▶ Play a video along with a CPU-intensive application
 - ▶ CPU-intensive application (e.g., compression)
 - ▶ `$ sudo apt-get install p7zip-full sysstat`
 - ▶ `$ 7z b 100 -md16`
 - ▶ Video (with a media player)
 - ▶ `$ mpv --loop=inf video.mp4`
 - ▶ Choose a suitable video that can be played in your VM and consumes some CPU resources
- ▶ Apply real-time scheduling for all threads and CPU-intensive threads
 - ▶ **Example:** `$ sudo chrt --all-tasks --deadline --sched-runtime 700000 --sched-deadline 10000000 --sched-period 10000000 -p 0 $(pgrep mpv)`
 - ▶ Observe and analysis dropped frames in different configurations

```
A0: [pulse] 44100Hz stereo 2ch float
VO: [sdl] 640x480 yuv420p
AV: 00:01:26 / 00:03:09 (45%) A-V: 0.394 Dropped: 12
```

Print Information for Real-Time Tasks

- ▶ Modify the kernel to print a message whenever scheduling parameters are changed (e.g., by chrt) for a task
- ▶ `$ sudo dmesg`

```
[10701.608517] SCHED_DEADLINE parameters set on pid 3792: runtime=700000, deadline=10000000, period=10000000  
[10701.608526] SCHED_DEADLINE parameters set on pid 3794: runtime=700000, deadline=10000000, period=10000000  
[10701.608532] SCHED_DEADLINE parameters set on pid 3795: runtime=700000, deadline=10000000, period=10000000
```

Hint

- ▶ The deadline scheduler is implemented in `kernel/sched/deadline.c`
- ▶ chrt uses `sched_setattr()` system call to assign the scheduler and scheduling parameters (runtime, deadline, period)
 - ▶ Parameters are passed via struct `sched_attr`
- ▶ The deadline scheduler maintains both assigned and runtime parameters
- ▶ The permission to assign a real-time scheduler
- ▶ Incremental kernel builds should be fast (a few minutes)

Project1 Requirements

- ▶ A PDF report within 2 pages and modified kernel source files
- ▶ Report should included: Linux Scheduling Policy Testing and Analysis and short description on kernel compilation steps (and problems, if any) and code tracing on deadline scheduler
 - ▶ For scheduling policies analysis, please compare the dropped frame among (1) different scheduling policies and (2) different parameters and priority combination
 - ▶ Describe WHAT you observe and WHY it behave like that
- ▶ Please show your names and student IDs in your report
- ▶ Describe work done by each member
- ▶ Be packed as one file named “RTS_PJ1_Team##_v#.zip”
 - RTS2024_PJ1_report.pdf
 - linux-5.15.71/kernel/sched/deadline.c

Submission Rules

- ▶ Project deadline: 2025/05/06 23:59
 - ▶ Delayed penalty: -20/Day
- ▶ Upload the zip file via NTU COOL
- ▶ **DO NOT COPY THE HOMEWORK**
 - ▶ Discussions among teams are encouraged, as long as properly credited

References

- ▶ Professional Linux® Kernel Architecture, Wolfgang Mauerer, Wiley Publishing, Inc.
- ▶ Deadline Task Scheduling
<https://www.kernel.org/doc/html/latest/scheduler/sched-deadline.html>